


Big Data Engineering With MapReduce and Hive


Agenda

hadoop




Hadoop Implementations

hadoop





Hadoop




Turn back the clock, the mainframe




Distributed Computing




Cloud Computing



If dogs were servers...




HDFS & MapReduce




Processing: 30 hours

HDFS & MapReduce



Processing: 15 hours

HDFS & MapReduce



Most cases, linear scaling of processing power

Number of Computers	Processing Time (hours)
1	30
2	15
3	10
4	7.5
5	6
6	5
7	4.26
8	3.75
9	3.33

Limitations with MapReduce

- ~200 lines of code to do anything
- Slow
- Troubleshooting multiple computers
- Good devs are scarce
- Expensive certifications

```
1 package org.apache.hadoop.examples;
2
3 import java.io.IOException;
4 import java.util.StringTokenizer;
5
6 import org.apache.hadoop.conf.Configuration;
7 import org.apache.hadoop.fs.Path;
8 import org.apache.hadoop.io.IntWritable;
9 import org.apache.hadoop.io.Text;
10 import org.apache.hadoop.mapreduce.Job;
11 import org.apache.hadoop.mapreduce.Mapper;
12 import org.apache.hadoop.mapreduce.Reducer;
13 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
14 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
15 import org.apache.hadoop.util.GenericOptionsParser;
16
17 public class WordCount {
18
19     public static class TokenizerMapper
20         extends Mapper<Object, Text, Text, IntWritable>{
21
22         private final static IntWritable one = new IntWritable(1);
23         private Text word = new Text();
24
25         public void map(Object key, Text value, Context context
26             ) throws IOException, InterruptedException {
27             StringTokenizer itr = new StringTokenizer(value.toString());
28             while (itr.hasMoreTokens()) {
29                 word.set(itr.nextToken());
30                 context.write(word, one);
31             }
32         }
33     }
```




Ambari: Cluster provisioning, management, and monitoring



Avro (Microsoft .NET Library for Avro): Data serialization for the Microsoft .NET environment



HBase: Non-relational database for very large tables



HDFS: Hadoop Distributed File System



Hive: SQL-like querying




Mahout: Machine learning

MapReduce and YARN: Distributed processing and resource management



Oozie: Workflow management



Pig: Simpler scripting for MapReduce transformations



Sqoop: Data import and export



Storm: Real-time processing of fast, large data streams



Zookeeper: Coordinates processes in distributed systems

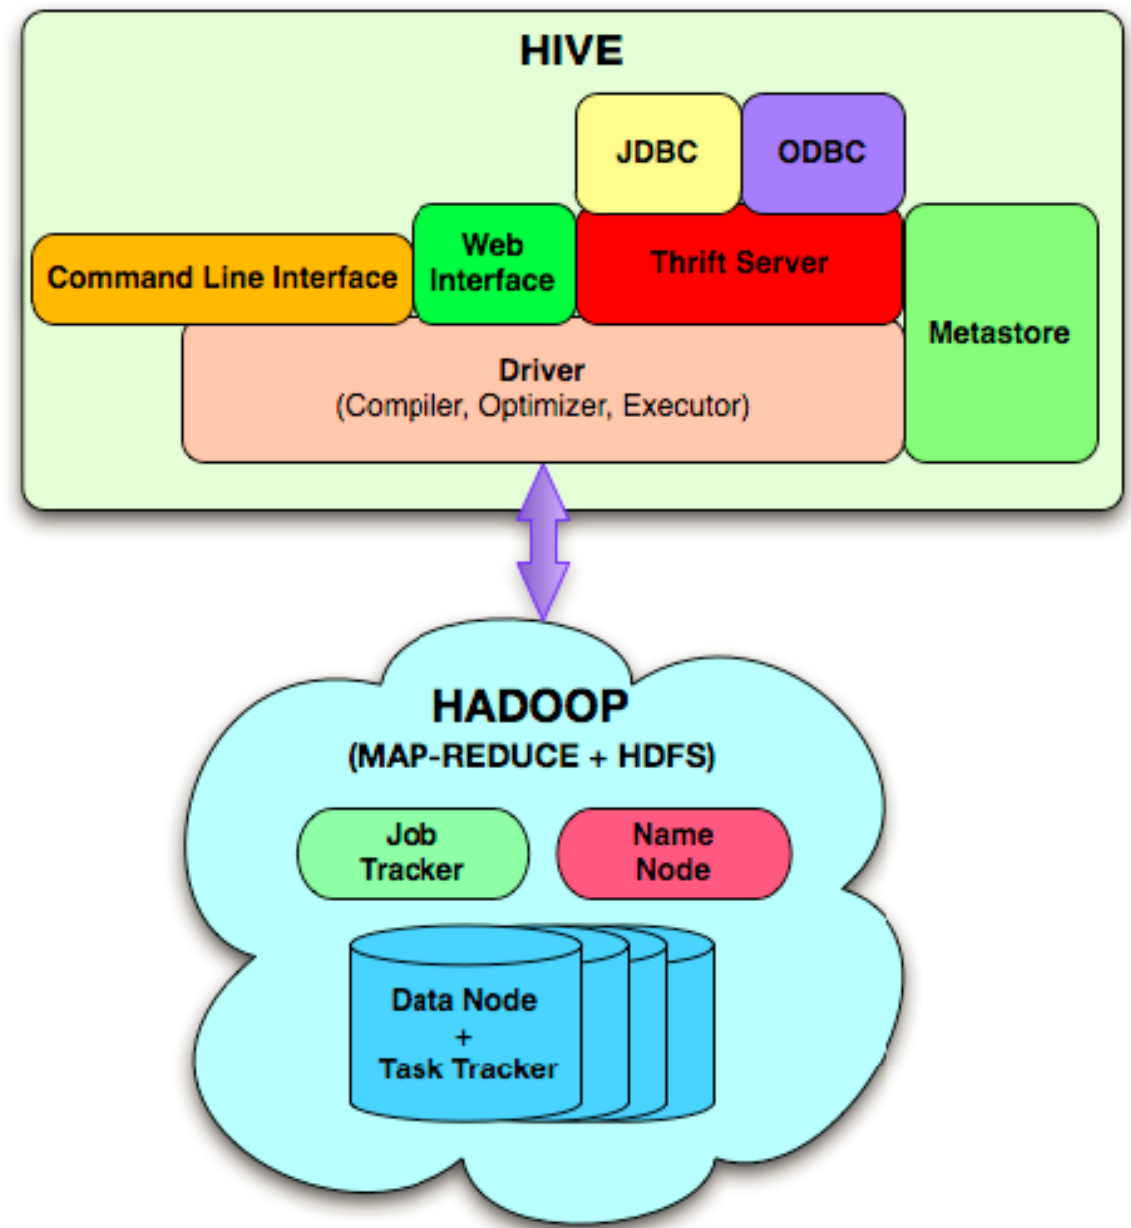



Fig. 1: Hive System Architecture

Hive Jobs



Hive Architecture





Data File



Unstructured
Data



Data File



Metadata File/DB



Structured
Data

Why Hive?

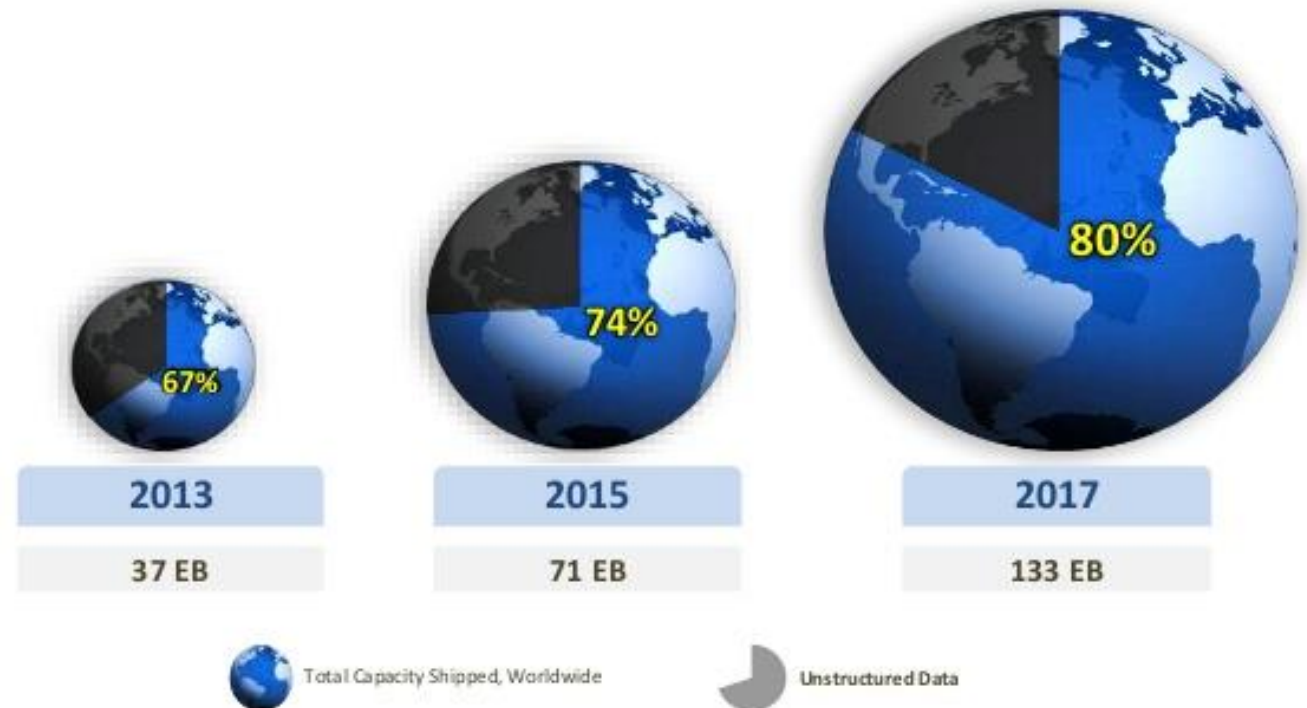


- SQL spoken here (HiveQL)
- ODBC driver
- BI Integration
- Supports only Structured Data

Limitations



Structured vs. Unstructured Data Growth





Source: IDC

Azure Blob Storage

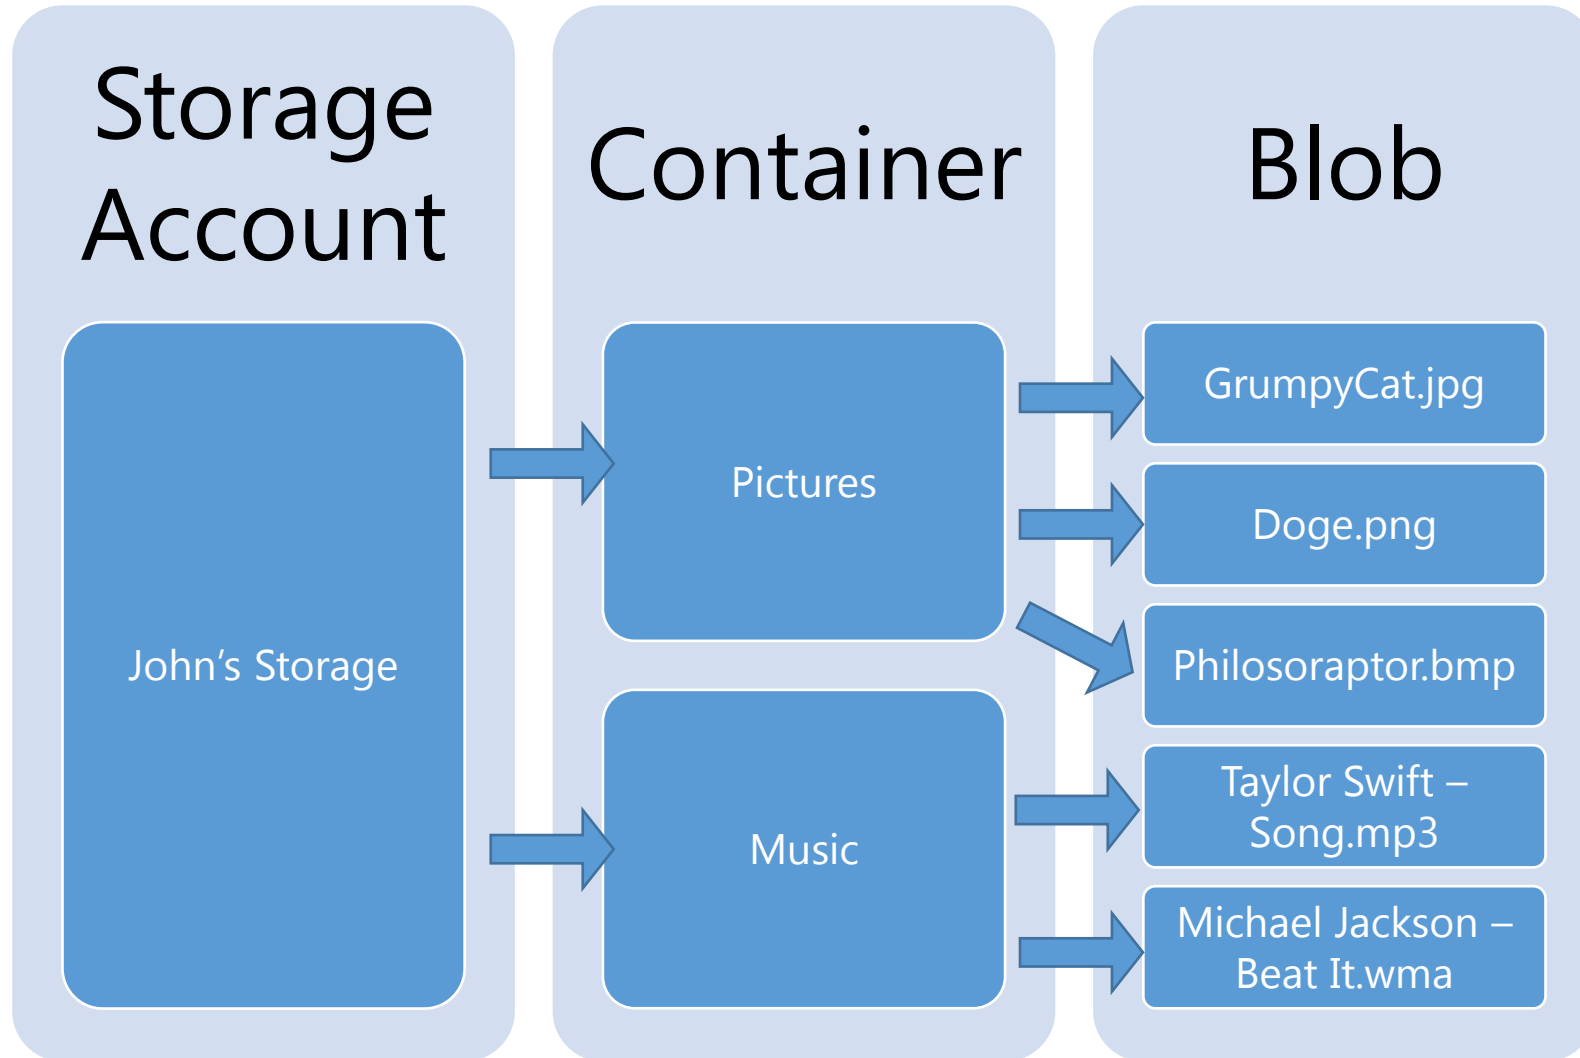


HDInsight




Blob Storage

Azure Blob Storage



When to Use Each




C#
Java
MapReduce

VS



Hive

VS



Pig

Questions?