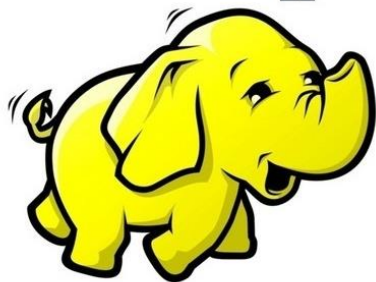


# Big Data Engineering With MapReduce and Hive

Data Science Dojo

# Agenda

*hadoop*




# Hadoop Implementations

***hadoop***




HDInsight



Amazon Elastic  
MapReduce

**data science dojo**  
unleash the data scientist in you


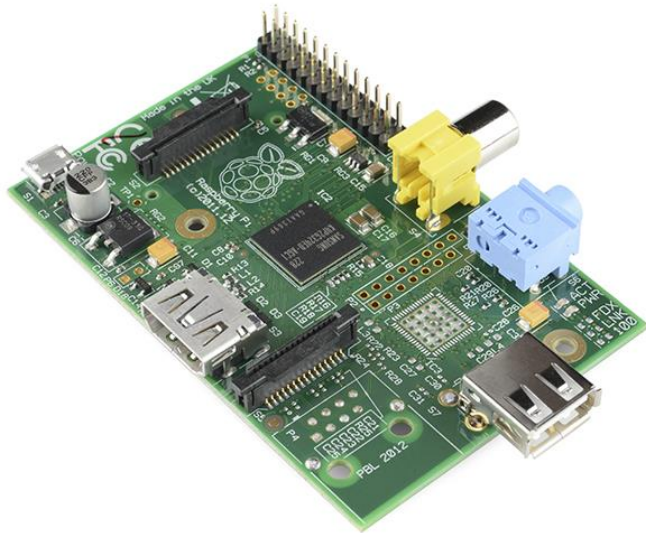
# Hadoop




# Turn Back The Clock, The Mainframe




# Distributed Computing



# Cloud Computing

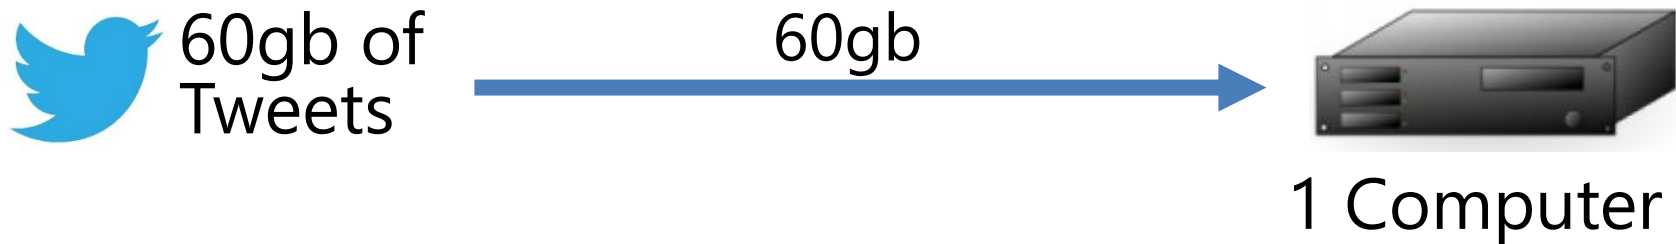


# If dogs were servers...






# HDFS & MapReduce



Processing: 30 hours

# HDFS & MapReduce



2 Computers

Processing: 15 hours

# HDFS & MapReduce



60 Gb of  
Tweets

20Gb

20Gb

20Gb



Processing: 10 hours


3 Computers

# Most Cases, Linear Scaling Of Processing Power




Number of Computers	Processing Time (hours)
1	30
2	15
3	10
4	7.5
5	6
6	5
7	4.26
8	3.75
9	3.33

# HDFS

## HDFS Partitioning



# HDFS Redundancy



# Limitations with MapReduce

- ~70 lines of code to do anything
- Slow
- Troubleshooting multiple computers
- Good devs are scarce
- Expensive certifications

```
1 package org.apache.hadoop.examples;
2
3 import java.io.IOException;
4 import java.util.StringTokenizer;
5
6 import org.apache.hadoop.conf.Configuration;
7 import org.apache.hadoop.fs.Path;
8 import org.apache.hadoop.io.IntWritable;
9 import org.apache.hadoop.io.Text;
10 import org.apache.hadoop.mapreduce.Job;
11 import org.apache.hadoop.mapreduce.Mapper;
12 import org.apache.hadoop.mapreduce.Reducer;
13 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
14 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
15 import org.apache.hadoop.util.GenericOptionsParser;
16
17 public class WordCount {
18
19     public static class TokenizerMapper
20         extends Mapper<Object, Text, Text, IntWritable>{
21
22         private final static IntWritable one = new IntWritable(1);
23         private Text word = new Text();
24
25         public void map(Object key, Text value, Context context
26             ) throws IOException, InterruptedException {
27             StringTokenizer itr = new StringTokenizer(value.toString());
28             while (itr.hasMoreTokens()) {
29                 word.set(itr.nextToken());
30                 context.write(word, one);
31             }
32         }
33     }
```



**Ambari:** Cluster provisioning, management, and monitoring



**Avro** (Microsoft .NET Library for Avro): Data serialization for the Microsoft .NET environment



**HBase:** Non-relational database for very large tables



**HDFS:** Hadoop Distributed File System



**Hive:** SQL-like querying



**Mahout:** Machine learning

**MapReduce and YARN:** Distributed processing and resource management



**Oozie:** Workflow management



**Pig:** Simpler scripting for MapReduce transformations



**Sqoop:** Data import and export




**Storm:** Real-time processing of fast, large data streams



**Zookeeper:** Coordinates processes in distributed systems






# Hive Jobs



# Hive Architecture





**Data File**



Unstructured  
Data



**Data File**



**Metadata File/DB**



Structured  
Data

# Semi Structured Data

## Self Describing Flat Files

- XML
- JSON
- CSV
- TSV

```
[  
  {  
    "created_at": "Thu May 07 18:06:23 +0000 2015",  
    "id": 596375540631646210,  
    "id_str": "596375540631646210",  
    "text": "Expert usable tips differently the press",  
    "source": "<a href=\\\"http://twitterfeed.com\\\" rel",  
    "truncated": 0,  
    "in_reply_to_status_id": null,  
    "in_reply_to_status_id_str": null,  
    "in_reply_to_user_id": null,  
    "in_reply_to_user_id_str": null,  
  }  
]
```


# Why Hive?



- SQL spoken here (HiveQL)
- ODBC driver
- BI Integration
- Supports only Structured Data

# Limitations

## Structured vs. Unstructured Data Growth





Source: IDC

# Azure Blob Storage



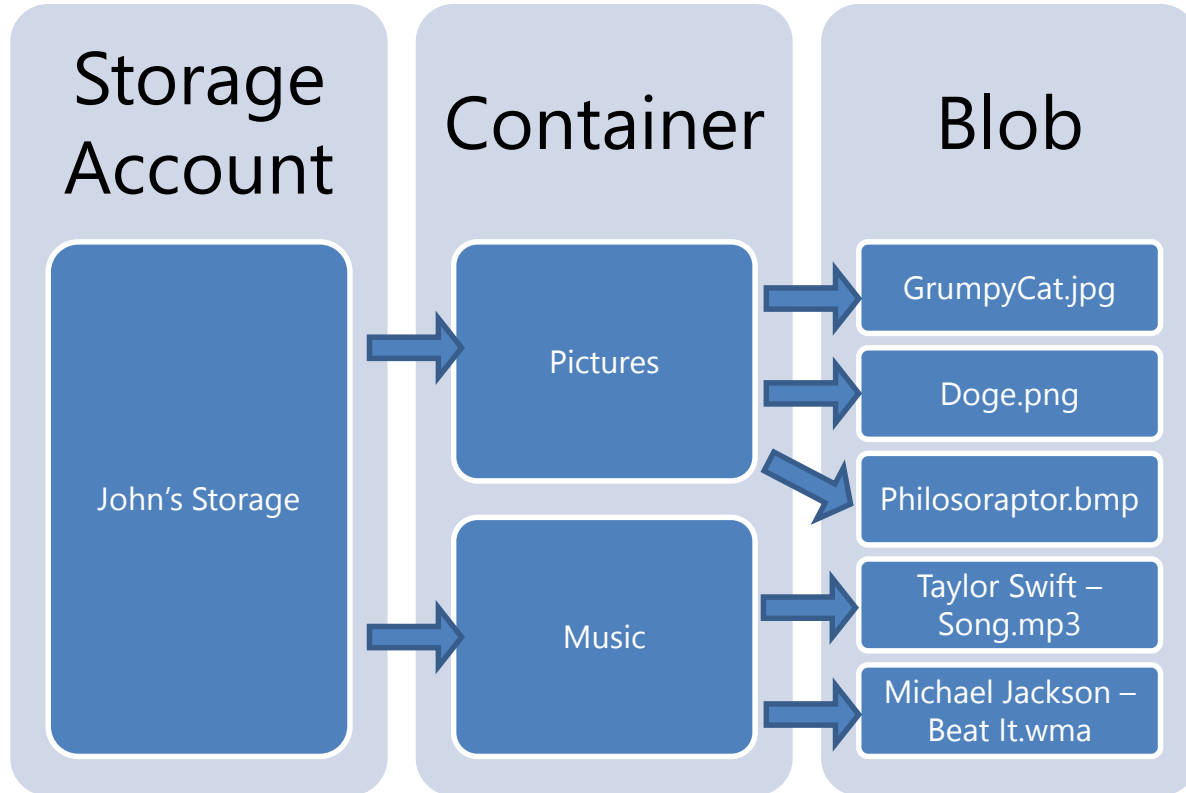
HDInsight




Blob Storage



# Azure Blob Storage



# When to Use Each




C#  
Java  
MapReduce

VS



Hive

VS



Pig


# MapReduce, via Playing Cards



Let's count the number of spades, clubs, hearts, and diamonds in a stack of cards, the way map reduce would.

- Each card represents a row of data
- Each suite represents an attribute of the data


# Using a 2 Data Node Cluster






# Mapping: Node Sorting


Node 1




Node 2



# Mapping: Node Shuffle, Data Transfer



# Mapping: Node Shuffle, Data Transfer





# Database = Normalization

select \* from **courses**;

id	name	teacher_id
10001	Computer Science 142	1234
10002	Computer Science 143	5678
10003	Computer Science 154	9012
10004	Informatics 100	1234

select \* from **students**;

id	name	email	password
123	Bart	bart@fox.com	bartman
404	Ralph	ralph@fox.com	catfood
456	Milhouse	milhouse@fox.com	fallout
789	Nelson	muntz@fox.com	haha!
888	Lisa	lisa@gmail.com	vegan

select \* from **grades**;

student_ids	course_id	grade
123	10001	B-
123	10002	C
456	10001	B+
888	10002	A+
888	10003	A+
404	10004	D+
456	10002	D-
404	10002	B
789	10003	D+

# Normalization, joining

```
select
    g.student_id,
    c.name as course,
    g.grade as grade
from grades as g
join courses c
    on g.course_id = c.id
join students s
    on g.student_id = s.id
;
```

student	course	grade
Bart	Computer Science 142	B-
Milhouse	Computer Science 142	B+
Bart	Computer Science 143	C
Lisa	Computer Science 143	A+
Milhouse	Computer Science 143	D-
Ralph	Computer Science 143	B
Lisa	Computer Science 154	A+
Nelson	Computer Science 154	D+
Ralph	Informatics 100	D+

# Data Warehouse = Denormalization


<b>student</b>	<b>course</b>	<b>grade</b>
Bart	Computer Science 142	B-
Milhouse	Computer Science 142	B+
Bart	Computer Science 143	C
Lisa	Computer Science 143	A+
Milhouse	Computer Science 143	D-
Ralph	Computer Science 143	B
Lisa	Computer Science 154	A+
Nelson	Computer Science 154	D+
Ralph	Informatics 100	D+

## Tables:


- Students Table
- Courses Table
- Roster Table

# Trade-Offs

## Normalization



## Denormalization




# Costs, Storage vs Processing

## Storage

US – N. Virginia	US – N. California	EU – Ireland
Standard On-Demand Instances	Linux/UNIX Usage	Windows Usage
Small (Default)	\$0.085 per hour	\$0.12 per hour
Large	\$0.34 per hour	\$0.48 per hour
Extra Large	\$0.68 per hour	\$0.96 per hour

## Processing

US – Standard	US –
Storage	
Tier	Pricing
First 50 TB / Month of Storage Used	\$0.150 per GB
Next 50 TB / Month of Storage Used	\$0.140 per GB
Next 400 TB /	\$0.130 per GB







- Distributed Machine Learning
- Installed into Hadoop & Spark
- R-like language Implementation

# Distributed Random Forest

HDFS  
Partitioning

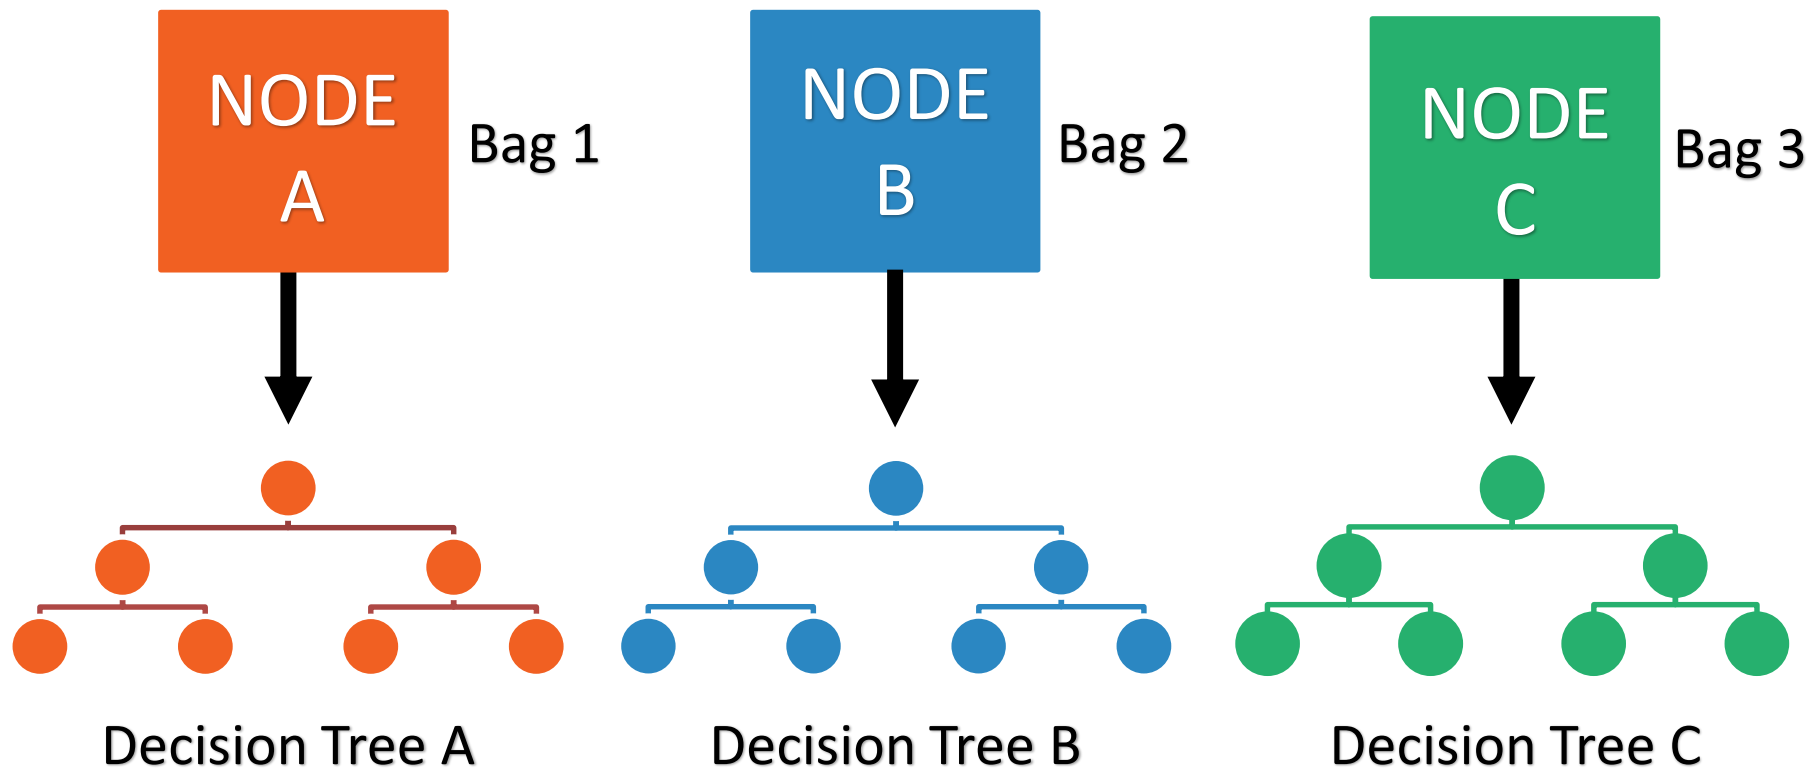


# Distributed Random Forest





# Distributed Random Forest



# Classification

	PC	Mahout	Spark
Logistic Regression - trained via SGD		<10m	
Naive Bayes			
Random Forest			
Hidden Markov Models			
Multilayer Perceptron			

Source: <https://mahout.apache.org/users/basics/algorithms.html>

# Recommendation Engines

	PC	Mahout	Spark
User-Based Collaborative Filtering			
Item-Based Collaborative Filtering			
Matrix Factorization with ALS			
Matrix Factorization with ALS on Implicit Feedback			
Weighted Matrix Factorization, SVD++			


Source: <https://mahout.apache.org/users/basics/algorithms.html>

# Clustering

	PC	Mahout	Spark
<b>k-Means Clustering</b>			
<b>Fuzzy k-Means</b>			
<b>Streaming k-Means</b>			
<b>Spectral Clustering</b>			

Source: <https://mahout.apache.org/users/basics/algorithms.html>





■ Hadoop  
■ Spark

In-Memory: 100x  
times faster than  
Hadoop



3x faster on 10x few machines

Datona GraySort Benchmark: Sort 100 TB of data


Previous World Record:

- Method: Hadoop
- Yahoo!
- 72 Minutes
- 2100 Nodes

2014:


- Method: Spark
- Databricks
- 23 Minutes
- 206 Nodes

## Activity in last 30 days




Patches

MapReduce Storm  
Yarn Spark



Lines Added

MapReduce Storm  
Yarn Spark



Lines Removed

MapReduce Storm  
Yarn Spark





Spark  
SQL


Spark  
Streaming

MLlib  
(machine  
learning)

GraphX  
(graph)

Apache Spark

# Technology adoption life cycle



Source: <http://carlosmartinezt.com/2010/06/technology-adoption-life-cycle/>

# QUESTIONS