

Linear Regression

Overview

- Difference between Regression and Classification
- Linear Regression
 - Motivating Example: Predicting Housing Prices
 - Hypothesis function
 - Cost Function
 - Gradient Descent
- Applying linear regression for hand digit recognition

Supervised Learning

- Data: $D = \{ d_1, d_2, d_3, \dots, d_n \}$ a set of n examples
 $d_i = \langle x_i, y_i \rangle$

x_i

- Input vector
- Independent variables
- Explanatory variables
- Features
- Predictors

y_i

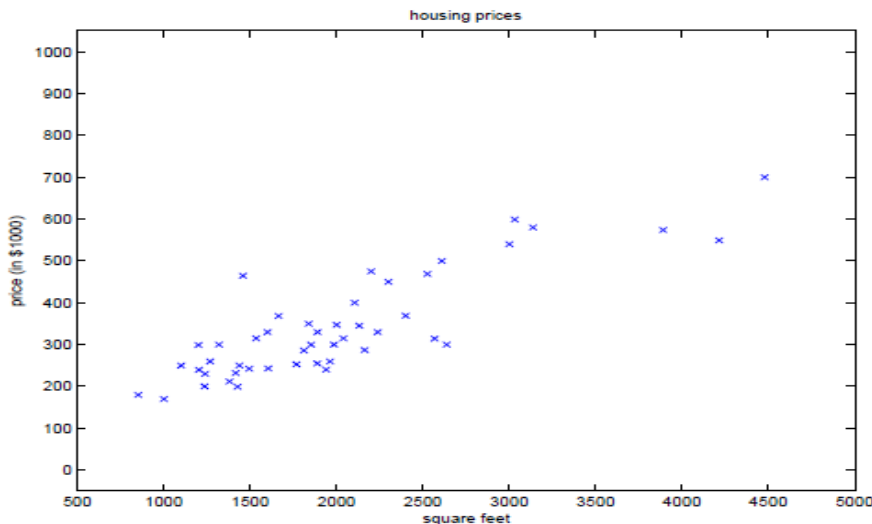
- Output scalar
- Dependent variables
- Response
- Outcome

Supervised Learning

- **Regression**: X discrete or continuous $\rightarrow Y$ is continuous
 - Example: Prices, Weight, Height, signal measurement, temperature etc.
- **Classification** : X discrete or continuous $\rightarrow Y$ is discrete
- **Objective**: learn the mapping of $f : X_i \rightarrow Y_i$

Predict Housing Prices

Living area (feet ²)	Price (1000\$)
2104	400
1600	330
2400	369
1416	232
3000	540
⋮	⋮




Can we learn to predict the price when the price of the house is not in the dataset?

Living Area = 3500 sq. ft.

Price = ?

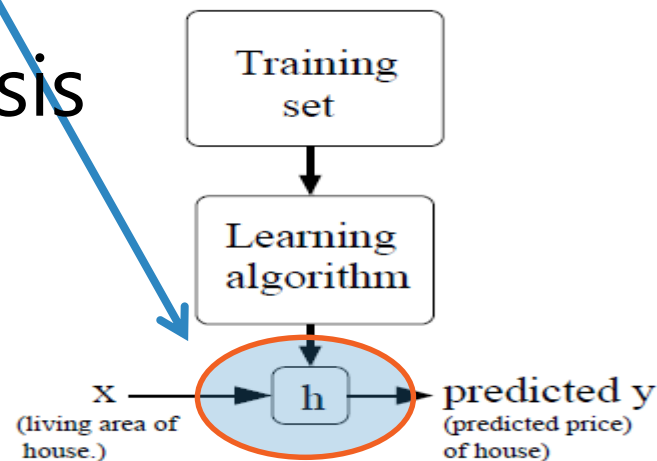
Predict Housing Prices

- Can we learn to predict the price when the price of the house is not in the dataset?
- Living Area = 3500 sq. ft. Price = ?
- Use the line that is somewhere in the middle
- How do we define somewhere in the middle?



Training Process In Linear Regression


- Objective: learn the mapping of $f : X_i \rightarrow Y_i$
- Finding h is the goal here
- h is known as the hypothesis



Predict Housing Prices

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Living area (feet ²)		Price (1000\$s)	
$x^{(1)}$	2104	400	$y^{(1)}$
$x^{(2)}$	1600	330	$y^{(2)}$
$x^{(3)}$	2400	369	$y^{(3)}$
$x^{(4)}$	1416	232	$y^{(4)}$
$x^{(5)}$	3000	540	$y^{(5)}$
\vdots	\vdots	\vdots	\vdots
$x^{(m)}$			$y^{(m)}$



m = Total number of training examples

θ_j = Parameters

Hypothesis Function-Linear Regression

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$


$h_{\theta}(x^i)$ is the prediction of the hypothesis h_{θ} for given input x^i
 y^i is the target values (what we would call labels in a classification problem)

We want the prediction $h_{\theta}(x^i)$ to be as close to the true label y^i as possible.

How do we do that?

Cost Function

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$





- In Plain English
 - find a linear line that is as close to the actual outputs as possible.
- Mathematically
 - Find linear function of X that minimizes the sum of squared residuals from Y.
 - a.k.a loss function

Hypothesis and Cost Function


Supposed:

$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$



The Cost Function



Hypothesis: $h_{\theta}(x) = \theta_0 + \theta_1 x$

Cost Function: $J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$

Parameters: θ_0, θ_1

Goal: minimize $J(\theta_0, \theta_1)$
 θ_0, θ_1

Gradient Descent Algorithm


- Goal : $\underset{\theta_0, \theta_1}{\text{minimize}} J(\theta_0, \theta_1)$
- Gradient descent starts with some initial θ and then performs an update for each value θ_j

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta).$$

- Repeat until θ converges

Gradient Descent Algorithm

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta).$$



Gradient Descent Algorithm

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta). \quad \text{For } j=0 \text{ and } j=1$$

Repeat until converge {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$

$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

}



Gradient Descent Algorithm

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta).$$


- α is known as the learning rate
- Each time the algorithm takes a step in the direction of the steepest, $J(\theta)$ decreases.
- α determines how quickly or slowly the algorithm will converge to a solution

Gradient Descent Algorithm

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta).$$



Digression - Gradient Descent Intuition



Visualize
animation in R

Multivariate Linear Regression

Predict Housing Prices – With Two Features


x_1^i Living area (feet ²)	x_2^i #bedrooms	y^i Price (1000\$s)
2104	3	400
1600	3	330
2400	3	369
1416	2	232
3000	4	540
\vdots	\vdots	\vdots

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

A linear function is just one of the choices to approximate the target variable.

θ is the space of linear functions mapping the space of input variables to the output/target variables

Predict Housing Prices – With Two Features



When we have two features, the linearity is in plane instead of line.

Algebraic Notation

The function approximating the target variable y is given by:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2$$

as

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x,$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \end{bmatrix}$$

$$\theta^T = [\theta_0 \ \theta_1 \ \theta_2]$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 1 \\ x_1 \\ x_2 \end{bmatrix}$$

$$\theta^T x = [\theta_0 \ \theta_1 \ \theta_2] \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix} = \sum_{i=0}^n \theta_i x_i$$

Batch Gradient Descent

For **ONE** training examples, we get the following update rule:

$$\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}.$$

What do we observe here about the magnitude of the update?

For **ALL** training examples, we get the following update rule:

Repeat until convergence {

$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j).$$

}

Stochastic Gradient Descent

- Consider the following algorithm:

```
Loop {  
    for i=1 to m, {  
         $\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}$     (for every  $j$ ).  
    }  
}
```

- This algorithm updates the parameters θ_j using each training example instead of all training examples.
- If the training set is big i.e., m is large, this technique converges quicker than batch gradient descent.
- Stochastic gradient descent may oscillate around the minimum of $J(\theta)$ and may not completely converge

Batch vs. Stochastic Gradient Descent

Batch Gradient Descent

Repeat until convergence {

$$\theta_j := \theta_j + \alpha \sum_{i=1}^m (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j).$$

}

- To update each parameter value, scan through the whole training data
- Converges to the minimum value slowly
- Preferred for small datasets

Stochastic Gradient Descent

Loop {

for i=1 to m, {

$$\theta_j := \theta_j + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)} \quad (\text{for every } j).$$

}


}

- Update the parameter values with one training example at a time
- Converges to the 'proximity' of minimum value fast but may keep oscillating near the minimum
- Preferred for large datasets

Example: Handwritten Digit Recognition



Extracting Features For Learning



$\{x_1, x_2, x_3, \dots, x_{256}, y = \text{'three'}\}$

- Each x_i corresponds to a feature value in the image
- y is a label of the training data; can be numeric or categorical, '3' or 'three'
- Each image is converted to row vectors and the appropriate learning algorithm is used
- Convention
 - x_i represents the i th feature in a training sample
 - y represents the label for the training sample

QUESTIONS