

K-Means Clustering

Unsupervised Learning

- Trying to find hidden structure in unlabeled data
- No error or reward signal to evaluate a potential solution
- Common techniques: K-Means clustering, hierarchical clustering, hidden Markov models, etc.
 - It has a long history, and used in almost every field, e.g., medicine, psychology, botany, sociology, biology, archeology, marketing, insurance, libraries, etc.

Unsupervised Learning

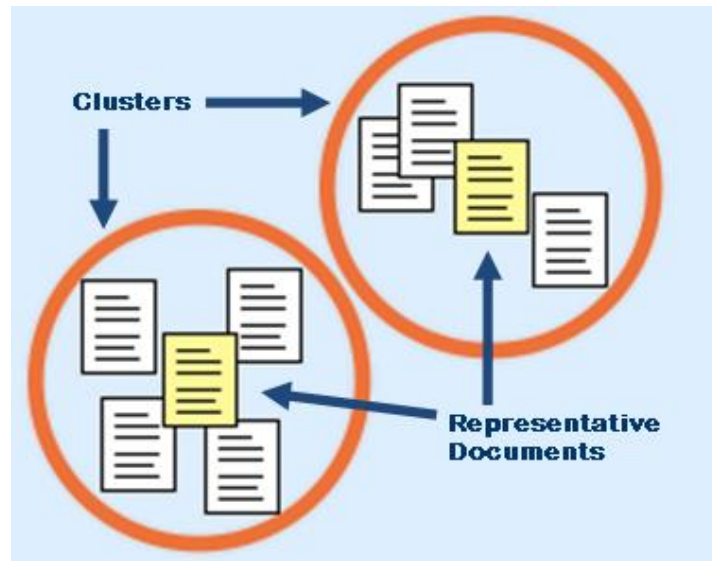
Example 1: Clothing size

- Tailor-made for each person is too expensive
- One-size-fits-all: does not work!
- Groups people of similar sizes together to make "small", "medium", and "large" t-shirts

Unsupervised Learning

Example 2: Text document organization


- To find groups of documents that are similar to each other based on the important terms appearing in them



Unsupervised Learning

Example 3: Target Marketing

- Subdivide market into distinct subsets of customers where any subset may conceivably be selected as a segment to be reached with a particular offer




K-Means Clustering

- Partitions data points into similarity clusters
- Only works for numeric data



Euclidean Distance


points in a two-dimensional space to determine intra- and inter-cluster similarity



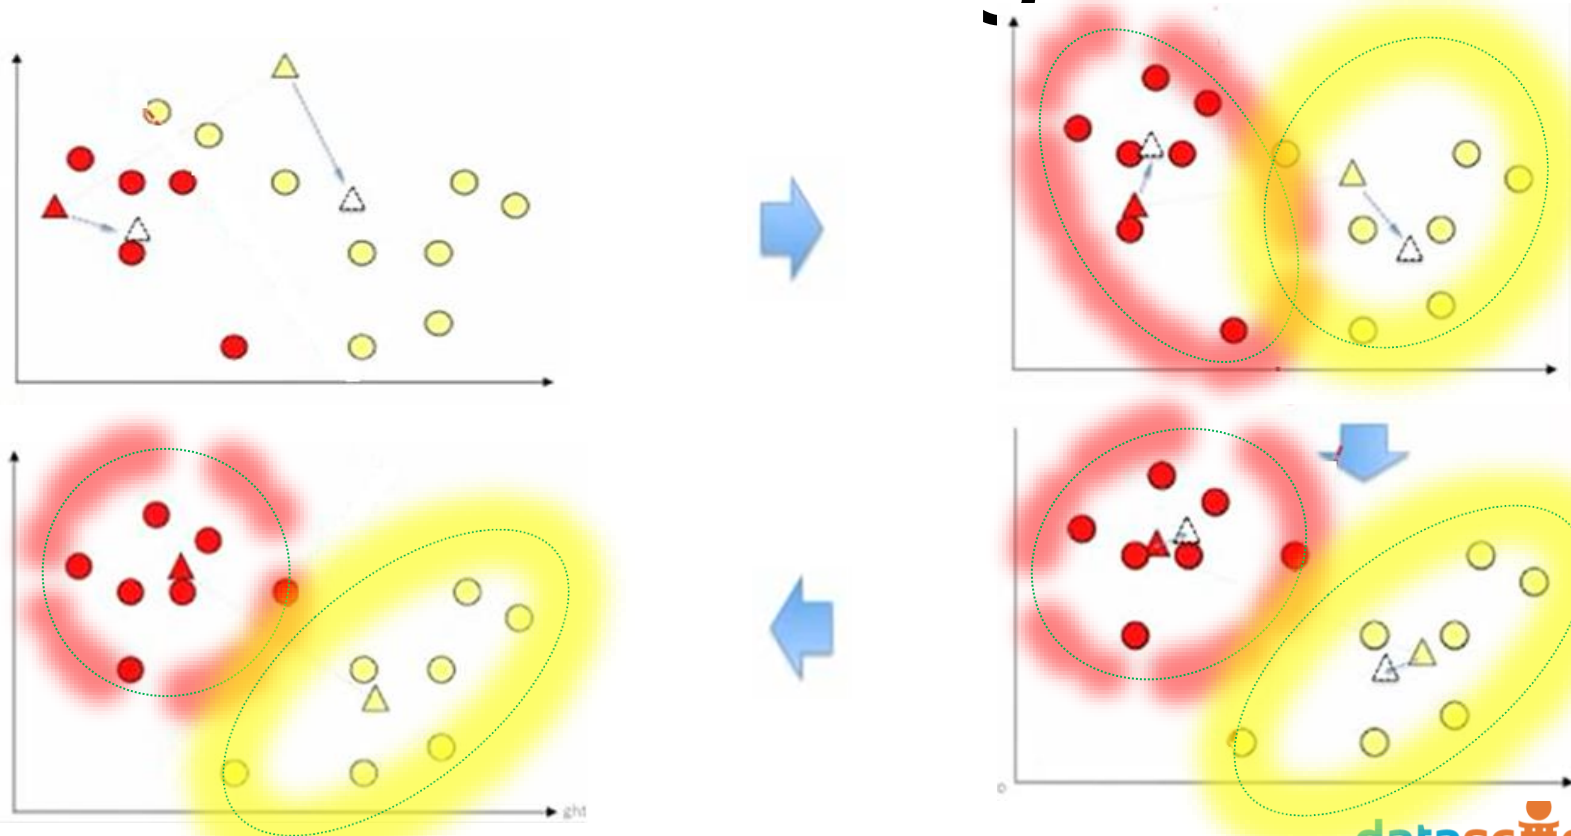
$$\sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$$

Intra-cluster distances
are minimized

Inter-cluster distances
are maximized



K-means Clustering




K-Means Clustering

- Minimizes aggregate intra-cluster distance
 - Measure squared distance from point to center of its cluster.

$$\sum_{j=1}^K \sum_{x \in g_j} D(c_j, x)^2$$

- Could converge to local minimum
 - Different starting points → very different results
 - Run many times with random starting points
- Nearby points may not be assigned to the same cluster



K-means Clustering

- Strengths
 - Simple: easy to understand and to implement
 - Efficient: linear time, minimal storage
- Weaknesses
 - Mean must be well defined
 - The user needs to specify K
 - Sensitive to outliers

How many clusters?


Rule of thumb

$$k \approx \frac{\sqrt{n}}{2}$$

n = number of data points

Elbow method

- percentage of variance explained as a function of the number of clusters
- choose a number of clusters so that adding another cluster doesn't give much better modeling of the data.




Other K Optimization Techniques

- Silhouette
- Calinsky criterion
- Bayesian Information Criterion
- Affinity propagation (AP) clustering
- Gap statistic

Example: Handwritten Digit Recognition



Extracting Features For Learning



$\{x_1, x_2, x_3, \dots, x_{256}, y = \text{'three'}\}$

- Each x_i corresponds to a feature value in the image
- y is a label of the training data; can be numeric or categorical, '3' or 'three'
- Each image is converted to row vectors and the appropriate learning algorithm is used
- Convention
 - x_i represents the i th feature in a training sample
 - y represents the label for the training sample

QUESTIONS